



Datenbanksysteme I

WS 2021/22

Torsten Grust, Christian Duta, Tim Fischer

Assignment #5

Submission Deadline: December 1, 2021 - 10:00

Exercise 1: Constraints

(14 Points)

Consider the following SQL DDL statement which creates a table to hold a company's employees:

```
CREATE TABLE employees (  
  employee_id  int,  
  lastname     text,  
  firstname    text,  
  address      text,  
  hire_date    date,  
  salary       salary,    -- monthly salary (in €)  
  emp_role     emp_role,  -- employee role  
  department_id int       -- identifier of employee's department  
);
```

Consider the following mini-world rules:

- i. No employee detail may be left unspecified.
 - ii. The **salary** of all employees must not be less than 1,473.33 € a month.
 - iii. An employee's role is either 'Manager', 'Developer', 'Accountant' or 'Secretary'.
 - iv. Managers hired after November 24, 2013 have a monthly **salary** of at most 17,679.96 €.
 - v. No two employees must share the same identifier (**employee_id**).
1. For each rule, use SQL DDL statements to define constraints to enforce the rule as well as one **INSERT**-statement that abides by the rule and one example that violates the rule. Create types for **salary** and **emp_role**.
 2. Please explain why it is impossible to enforce...
 - (a) ...rule **iv** using a domain constraint (**CREATE DOMAIN**),
 - (b) ...rule **v** in terms of a **CHECK** constraint.

Exercise 2: Defining Keys

(3 Points)

Please define the terms “candidate key”, “superkey” and “primary key” in your own words as precisely as possible.

Exercise 3: Identifying Keys**(5 Points)**

Based on the following instance of a table $T(A, B, C, D)$:

A	B	C	D
1	1	9	11
1	2	8	12
2	2	7	13
3	4	4	11
4	4	5	12
5	5	6	13

1. List all possible *candidate keys*.
2. List all superkeys of this table.

Exercise 4: Using Keys**(8 Points)**

Consider the following schema definition and constraints for table r :

```
CREATE TABLE r (a int, b varchar(9999), c int, d int, e text);
```

```
ALTER TABLE r ALTER COLUMN a SET NOT NULL;
```

```
ALTER TABLE r ALTER COLUMN b SET NOT NULL;
```

```
ALTER TABLE r ALTER COLUMN c SET NOT NULL;
```

```
ALTER TABLE r ADD UNIQUE (a, c);
```

```
ALTER TABLE r ADD UNIQUE (b);
```

```
ALTER TABLE r ADD UNIQUE (d);
```

1. Your task is to choose a primary key for r under the assumption that neither a nor c are unique on their own. Which (combinations of) columns are eligible for primary key? Which primary key would you choose? Why?
2. Extend the schema definition to declare the primary key you chose.
3. For one reason or another you are not satisfied with your choice of primary key in 4.1 above. Your co-worker thus proposes to add an **artificial primary key column** id to table r .

```
-- PostgreSQL assigns <table name>_pkey as the primary key constraint name, thus:
ALTER TABLE r DROP CONSTRAINT r_pkey;
```

```
ALTER TABLE r ADD COLUMN id int;
```

```
ALTER TABLE r ADD PRIMARY KEY (id);
```

What advantages and disadvantages does the new column id bring with it? Discuss briefly.

4. Consider the following query:

```
SELECT DISTINCT v.a, v.b, v.e FROM r AS v;
```

Is it necessary to use `DISTINCT` here? Justify your answer.