

INTRODUCTION TO RELATIONAL DATABASE SYSTEMS

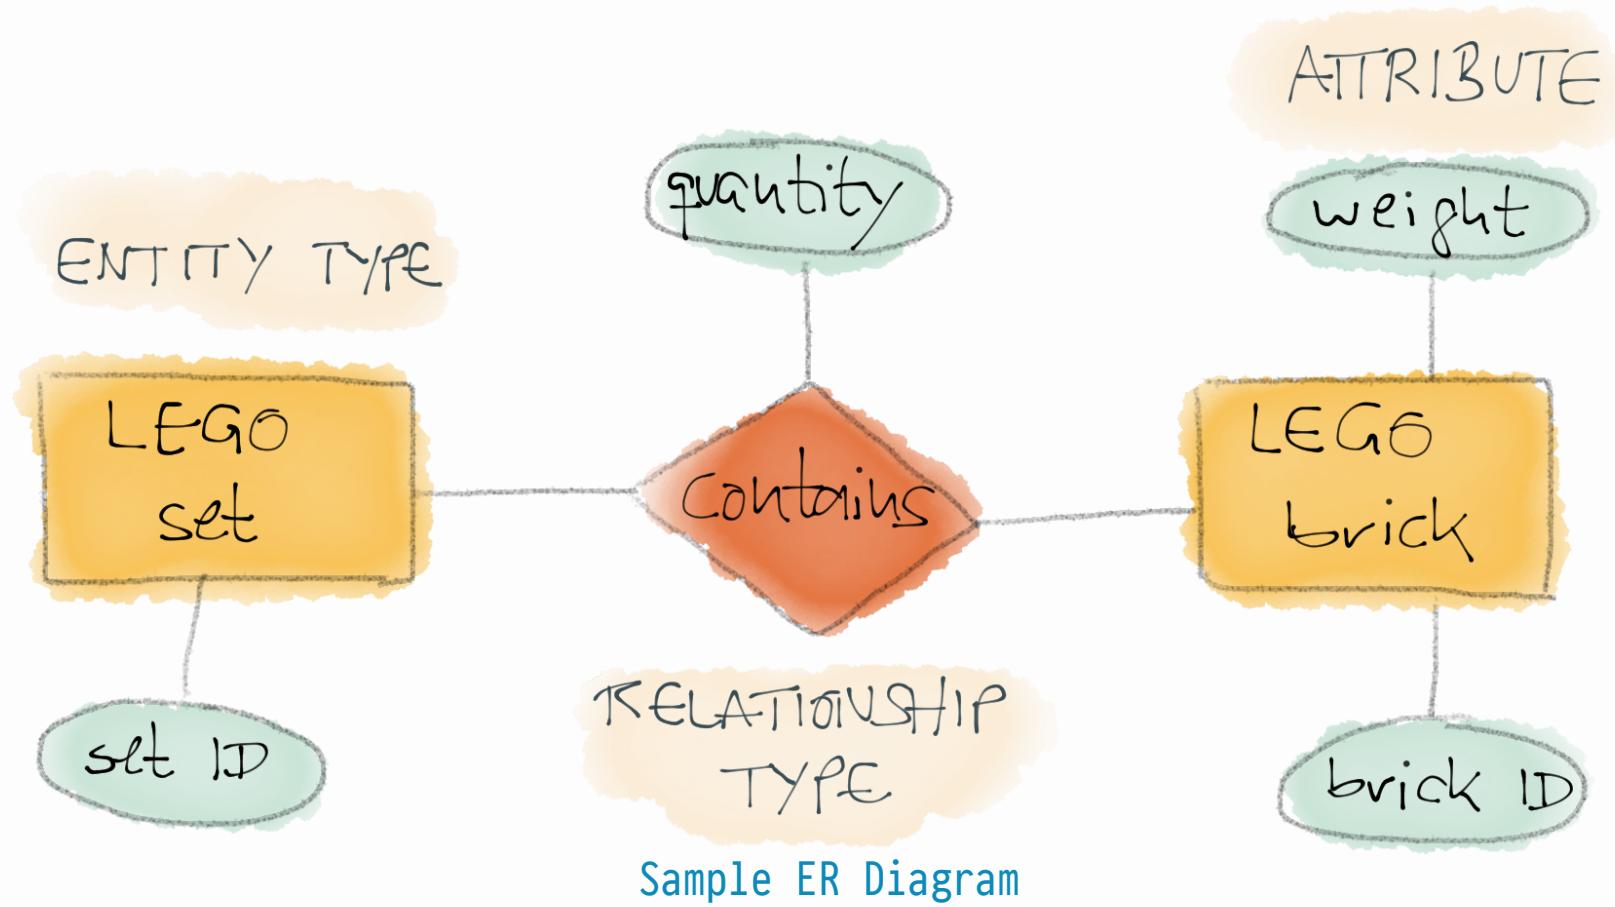
DATENBANKSYSTEME 1 (INF 3131)

Torsten Grust
Universität Tübingen
Winter 2021/22

THE ENTITY-RELATIONSHIP MODEL

- The **Entity-Relationship Model (ER model)** has been defined to concisely describe mini-worlds:
 - ER describes objects (**entities**) and their **relationships**. Objects and relationships carry **attributes** that characterize them further.
 - ER comes with a **graphical notation (ER diagrams)** that helps to establish a quick overview of complex mini-worlds. Also a great way to communicate models to domain experts/DB non-experts/future DB users.
- There are *no* ER model DBMSs. ER is also known as a **semantic data model**. Instead, **ER diagrams are translated** into the constructs of a concrete data model (e.g., relational or JSON) to generate a database schema.
- While the ER model/diagram is being designed, the focus exclusively remains on mini-world aspects. Details of the implementing data model — e.g., types, constraints — only come into play when the (automatic) translation is performed.

ER: DIAGRAMS



ER CONSTRUCTS: ENTITIES, ENTITY TYPES

Entity and Entity Types

Entities represent relevant mini-world objects to be held in the database. Examples: LEGO sets, bricks, earthquakes, calendars.

Each entity has a specific entity type that defines its attributes.

The mini-world can contain only a finite number of objects. Entities are distinguishable from one another, i.e., entities possess some form of identity.

- Notes:
 1. Entities do not have to correspond to objects of physical existence but may also represent conceptual objects like, for example, *vacations*.
 2. Entity identity is inherent, *not* based on (attribute) values.
Upon translation, however, entity identity may often be implemented in terms of attribute values (e.g., through LEGO brick IDs, travel agency booking numbers).

ER CONSTRUCTS: RELATIONSHIPS, RELATIONSHIP TYPES

Relationship and Relationship Type

Relationships establish a connection between two entities (binary relation). One entity may participate in multiple relationships (or none). See **cardinalities** below.

Each relationship has a specific **relationship type** that defines its attributes and the two participating entity types.

Examples:

“A LEGO set <entity type> **contains** <relationship type> one or more LEGO bricks <entity type>.” “An earthquake <entity type> **occurs in** <relationship type> one given geographical region <entity type>.”

- Relationship types may be viewed from the angle of both participating entity types: “A LEGO set *contains* one or more LEGO bricks.” — “A LEGO brick is *contained in* one or more LEGO sets.”

ER CONSTRUCTS: ATTRIBUTES

Attributes

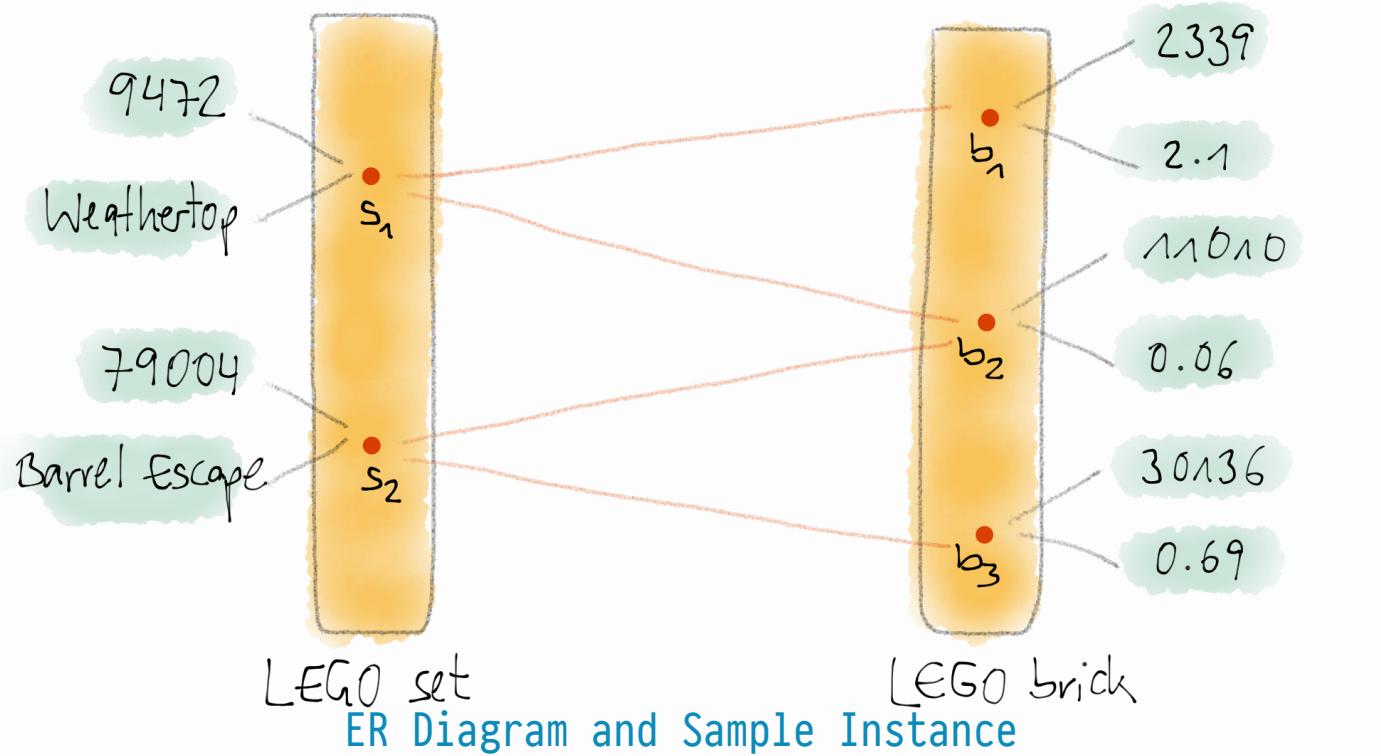
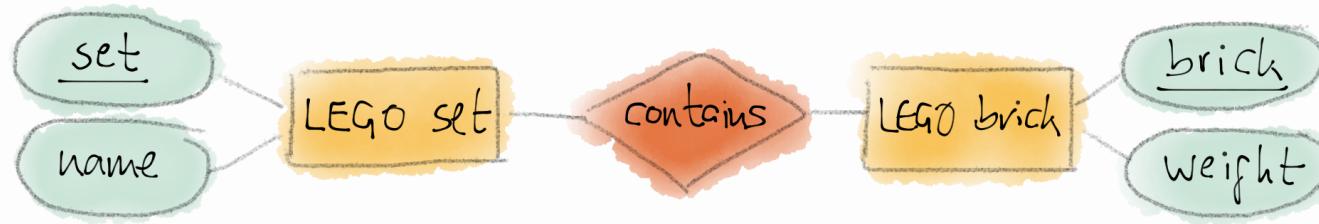
Entity and relationship types may carry **attributes** to represent properties of entities and relationships.

Attribute values may be of arbitrary data types (typically simple/atomic, e.g., string, number, date...). ER is first-order: attributes may *not* have values of type entity or relationship. Upon translation, the target data model deals with the representation of attribute values.

Selected attributes of an entity type may form a **key** (those attributes are underlined in ER diagrams): no two entities of the same entity type may feature the same key values.

- Key attributes exist *in addition* to entity identity and are helpful while translating an ER model into a target data model with value-based keys (e.g., relational).

SAMPLE ER INSTANCE (MINI-WORLD STATE)



ER DIAGRAM SEMANTICS

ER Diagram Semantics

The **ER Diagram Semantics** $\mathbb{E}\mathbb{R}$ interpret the symbols of an ER diagram by defining

1. a finite set $\mathbb{E}\mathbb{R}(e)$ (of entities) for every **entity type** e ,
2. a mapping $\mathbb{E}\mathbb{R}(a) : \mathbb{E}\mathbb{R}(e) \rightarrow \mathbb{V}(a)$ for every **attribute a of an entity type e** (if a is a key attribute, the mapping is injective),
3. a binary relation $\mathbb{E}\mathbb{R}(r) \subseteq \mathbb{E}\mathbb{R}(e_1) \times \mathbb{E}\mathbb{R}(e_2)$ for every **relationship type r** between entity types e_1 and e_2 ,
4. a mapping $\mathbb{E}\mathbb{R}(a) : \mathbb{E}\mathbb{R}(r) \rightarrow \mathbb{V}(a)$ for every **attribute a of a relationship type r** .

- Recall that $\mathbb{V}(a)$ denotes the set of admissible values for attribute a .

ER DIAGRAM SEMANTICS

- Example: The ER diagram semantics associated with the ER diagram and instance shown earlier defines:

1. $\text{ER}(\text{LEGO set}) = \{s_1, s_2\}$

$\text{ER}(\text{LEGO brick}) = \{b_1, b_2, b_3\}$

2. $\text{ER}(\text{set}) = f_1, \text{ER}(\text{name}) = f_2, \text{ER}(\text{brick}) = f_3, \text{ER}(\text{weight}) = f_4$
with

$$f_1(s_1) = 9472$$

$$f_2(s_1) = \text{'Weathertop'}$$

$$f_3(b_1) = 2339$$

$$f_4(b_1) = 2.1$$

$$f_1(s_2) = 79004$$

$$f_2(s_2) = \text{'Barrel Escape'}$$

$$f_3(b_2) = 11010$$

$$f_4(b_2) = 0.06$$

$$f_3(b_3) = 30136$$

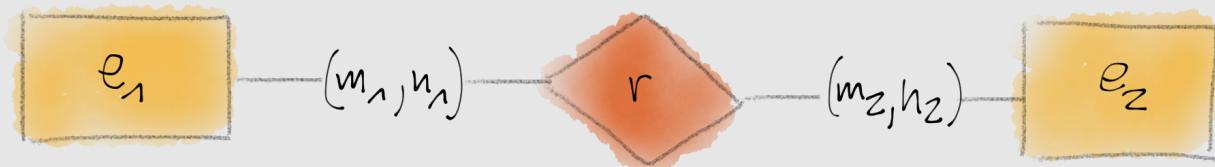
$$f_4(b_3) = 0.69$$

3. $\text{ER}(\text{contains}) = \{(s_1, b_1), (s_1, b_2), (s_2, b_2), (s_2, b_3)\}$

ER RELATIONSHIP CARDINALITIES

- In general, there is no restriction on **how often** a given entity **participates** in a relationship r .
- Specific **application semantics**, however, may dictate that participation is optional, mandatory, or that a minimum and/or maximum number of participations is required.

Relationship Cardinality (min/max Notation)



$$\begin{aligned} \forall e \in \text{ER}(e_1): m_1 &\leq | \{ (x,y) \in \text{ER}(r) \mid x = e \} | \leq n_1 \\ \forall e \in \text{ER}(e_2): m_2 &\leq | \{ (x,y) \in \text{ER}(r) \mid y = e \} | \leq n_2 \end{aligned}$$

- Notation: $n_i = *$ is interpreted as $n_i = \infty$.

ER DIAGRAM EXAMPLES

- Sketch ER diagrams to model the following mini-worlds. Identify entities and relationships, then restrict relationship cardinalities as required by the application:
 1. “A *man* can be *married to* at most one *woman* and vice versa.”
 2. “An *airport* lies in exactly one *country*. A *country* may have arbitrarily many *airports* (and maybe none at all).”
 3. “*Orders* are placed by *customers*. Some *customers* might not have placed an *order* yet.”
 4. “An *order* can contain several *products*.”
 5. “In the Euclidean 2d plane, a *line segment* connects two *points*. A simple closed polygon is formed by a list of three or more *line segments*. The interior of a polygon is colored.”

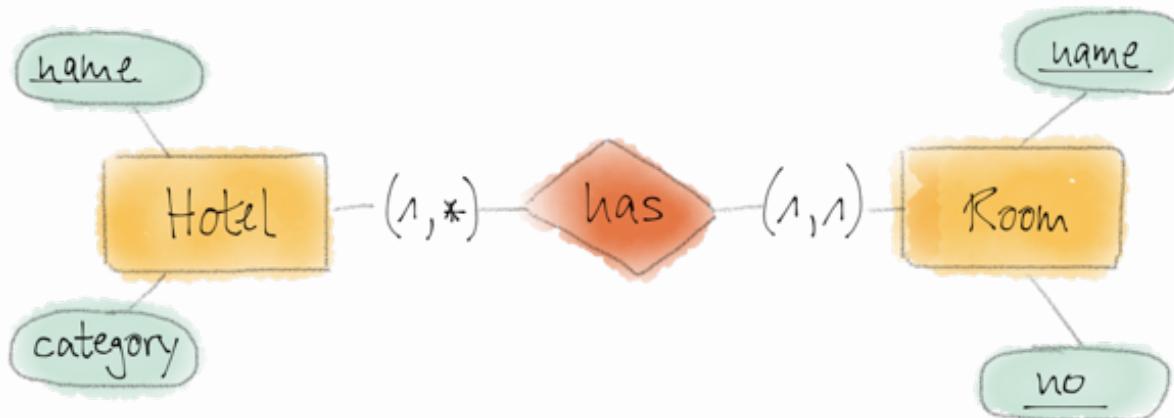
ER: NOTES ON RELATIONSHIP CARDINALITIES

- Cardinality $(0,*)$ represents an unrestricted relationship. A cardinality restriction (m_1, n_1) is **weaker** than (m_2, n_2) if $m_1 \leq m_2$ and $n_2 \leq n_1$.
- Relationship cardinalities denote **constraints** and have to be translated as such into the target data model.
 - If the target is relational, the important cardinalities $(0,1)$, $(1,1)$, and $(0,*)$ can be directly enforced by relational constraints, e.g., through **NOT NULL** or keys.
- ⚠ General cardinality constraints (m,n) may not translate (directly), however.
- Conventionally, relationships are categorized by their **maximum cardinalities** on both sides:

$(m_1, n_1) \leftrightharpoons (m_2, n_2)$	Relationship Category	
$n_1 = 1, n_2 = 1$	one-to-one (1:1)	
$n_1 = *, n_2 = 1$	one-to-many (1:N)	A
$n_1 = 1, n_2 = *$	many-to-one (N:1)	A
$n_1 = *, n_2 = *$	many-to-many (M:N)	

EXTENDED ER (EER): WEAK ENTITIES

- Common scenario: use an entity type to represent a **detail** of a superordinate **master** entity type. Without the master, the detail would not exist.
- Example:



- “One hotel has many rooms.” (General: “One master has many exclusive details.”)
- Detail entity type **Room** needs to form a **composite key** (incorporating the key of master entity type **Hotel**) for full identification.
- Implicit constraint: if detail entity e_2 is in relationship with master entity e_1 , both agree in the master’s key (here: attribute **name**).

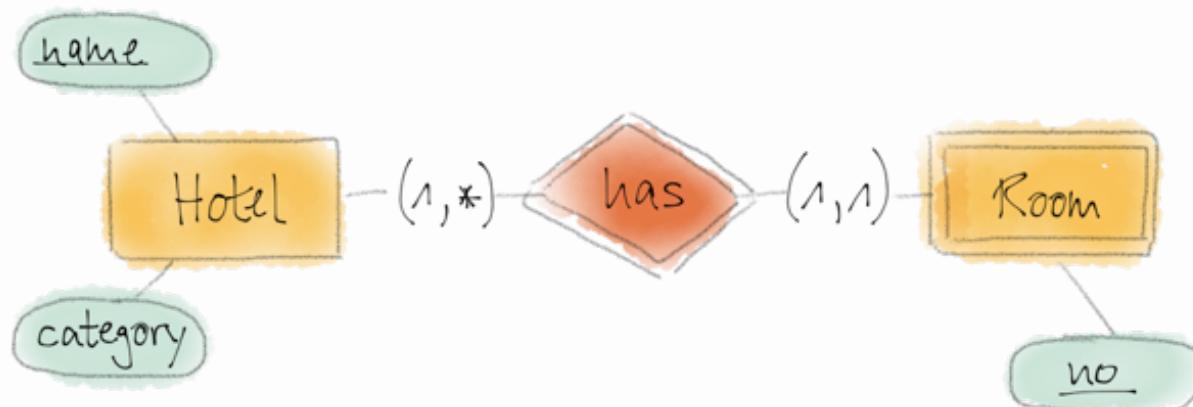
EER: WEAK ENTITIES

Weak Entity Type

In ER diagrams, a **weak entity type** (detail) and its relationship with the master entity type are drawn using **double-stroked lines**.

The weak entity type **implicitly inherits the master key's attributes** (and adds its own key attributes to form a composite key). The **existence of weak entities depends on their master entity** (affects translation to the target data model).

- Example (continued):



EER: WEAK ENTITIES

- Examples of master-detail scenarios in mini-worlds (identify existence dependencies and composite keys):
 1. “An **invoice** has a number of positions.”
 2. “A section in a book is identified by a **chapter** and section title.”
 3. “A web page URI is composed of a **web server** DNS address and a path on that server.”
- Develop an ER diagram to model quizzes (multiple choice tests):

*"Each quiz is identified by a title, each question within a quiz is numbered, and each possible answer to a given question is referenced by a letter.
For each question and answer, the associated text is stored.
Answers are classified into correct and incorrect ones."*

(What is the complete key for each of the occurring entity types?)

EER: INHERITANCE

Inheritance (Entity Sub/Supertypes)

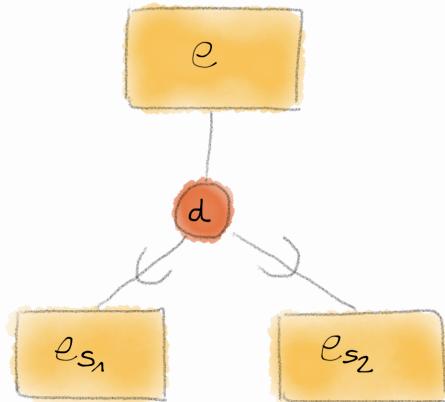
Entity type e_s is a **subtype** of entity type e if every entity of subtype e_s indeed also is an entity of **supertype** e (i.e., $\mathbb{ER}(e_s) \subseteq \mathbb{ER}(e)$). (**Specialization**)

Subtype e_s **inherits all attributes and relationships** from e . On top of these e_s may add further attributes and relationships.

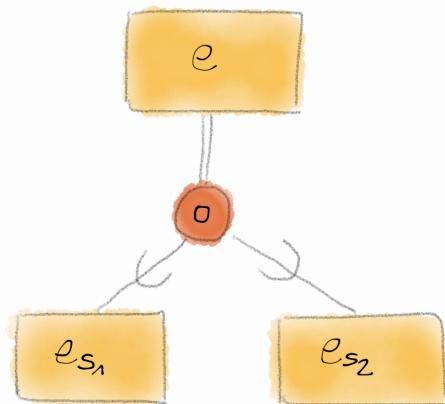
- Supertype e may have multiple subtypes e_{s1}, \dots, e_{sn} . Consider the following inheritance scenarios:
 - 1. Disjoint** (vs. overlap):
One entity may belong to a *single* subtype only, i.e., $\forall i \neq j: \mathbb{ER}(e_{si}) \cap \mathbb{ER}(e_{sj}) = \emptyset$.
 - 2. Total** (vs. partial):
Every entity *must* be a member of a subtype, i.e., $\mathbb{ER}(e) = \mathbb{ER}(e_{s1}) \cup \dots \cup \mathbb{ER}(e_{sn})$ (cf. with **abstract** classes in OOPs).

EER: INHERITANCE

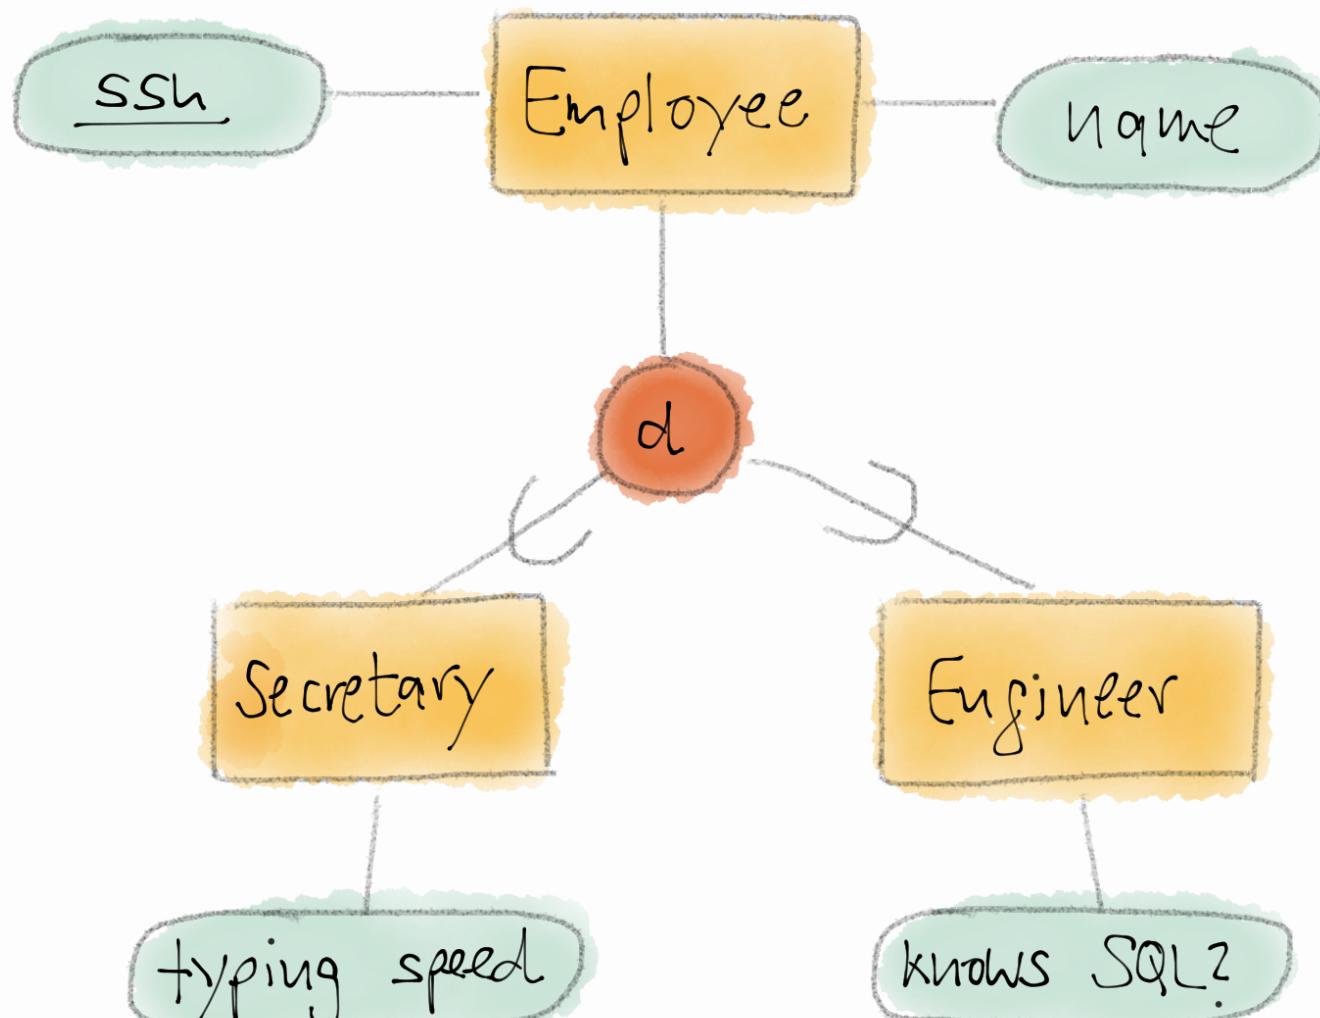
1. EER diagram: partial inheritance ($\text{--}\circlearrowleft$), disjoint subtypes (d in \circlearrowleft):



2. EER diagram: total inheritance ($=\circlearrowleft$), overlapping subtypes (o in \circlearrowleft):

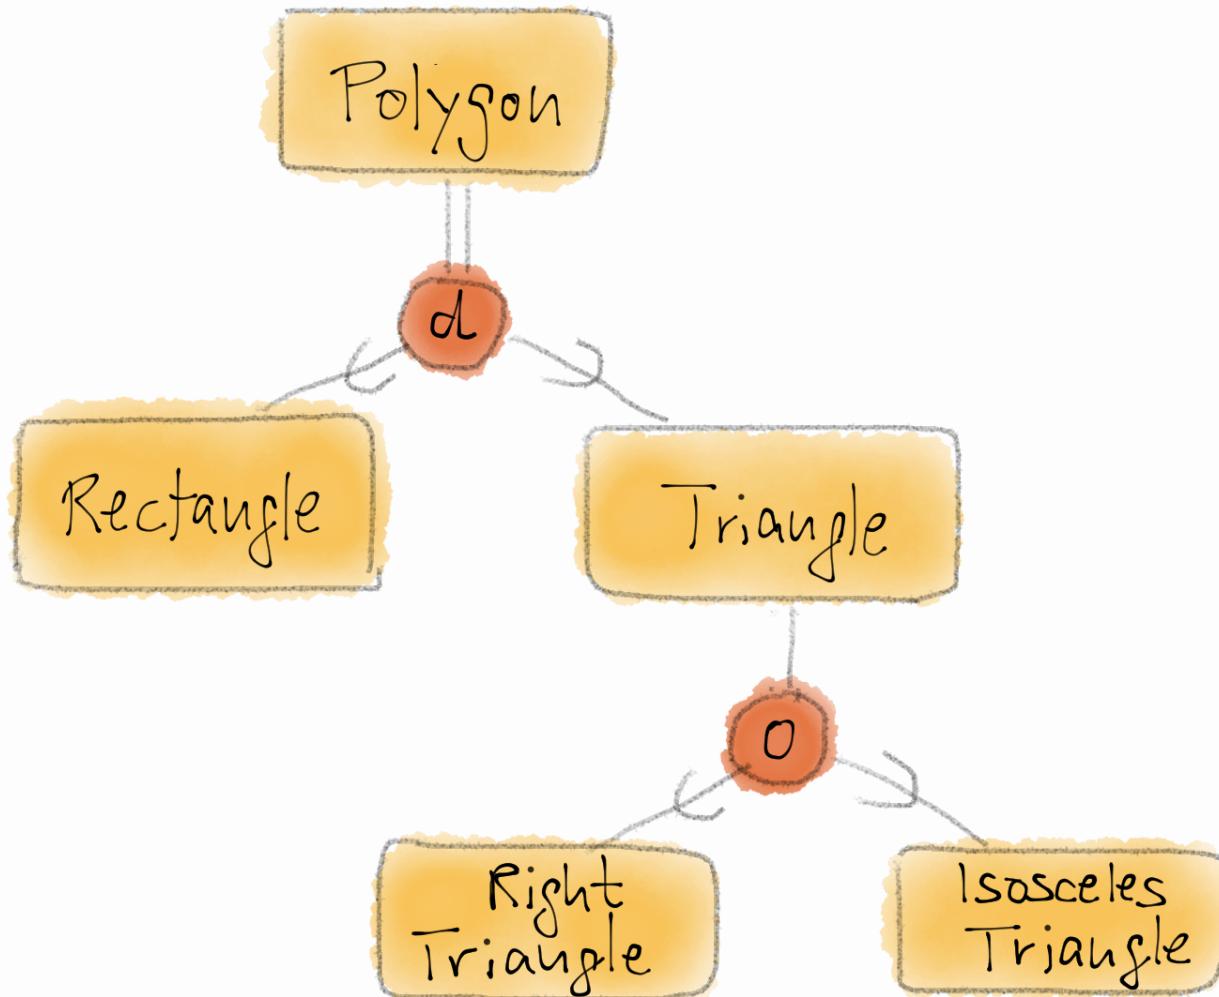


EER: INHERITANCE (EXAMPLES)



EER diagram example

EER: INHERITANCE (EXAMPLES)



EER diagram example (specialization hierarchy)

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #1: Entity Types

- Transforming entity type e :

1. **CREATE TABLE** named $t = \text{plural}(e)$

(each row of the table will hold one instance of type e).

- Pluralization (“*class*” → “*classes*”, “*story*” → “*stories*”, “*person*” → “*people*”):

(e.g., Damian Conway, Monash U, “A pluralization algorithm for English”,

<http://www.csse.monash.edu.au/~damian/papers/HTML/Plurals.html>)

- Python: `plural()` (in package/module `inflect`)

- Ruby: `String#pluralize` (see *Ruby on Rails’ ActiveSupport*)

2. Each **attribute** a of entity type e becomes a **column** a of table t (map attribute data type to SQL data type with appropriate value domain). If any, list attributes marked as key first.

MAPPING EER TO RELATIONAL TARGET (SQL)

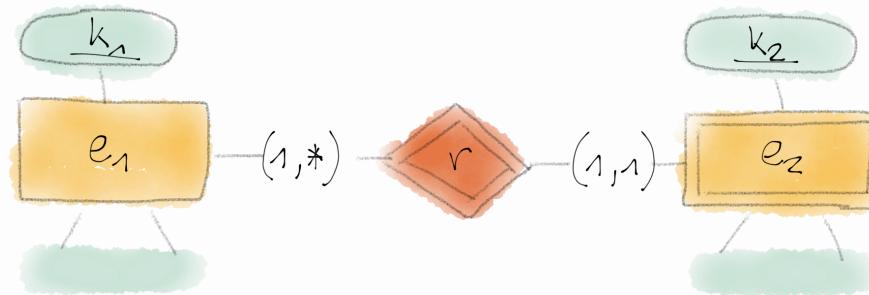
Step #1b: Entity Type Keys

- Transforming entity type e (target table t):
 1. If present, a **value-based primary key** of the entity type becomes the **primary key of table t**
`(ALTER TABLE t ADD PRIMARY KEY (...))`
 - If e 's key is composite, so will be the relational key.
 2. Otherwise, implement the **inherent entity identity** in terms of an **artificial key column**, say `_id_`, of auto-incrementing **integer** type:

```
CREATE TABLE  $t$  (...);  
-- establish artificial key column to implement entity identity  
ALTER TABLE  $t$  ADD COLUMN _id_ integer GENERATED ALWAYS AS IDENTITY;  
ALTER TABLE  $t$  ADD PRIMARY KEY (_id_);
```

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #2: Weak Entity Types

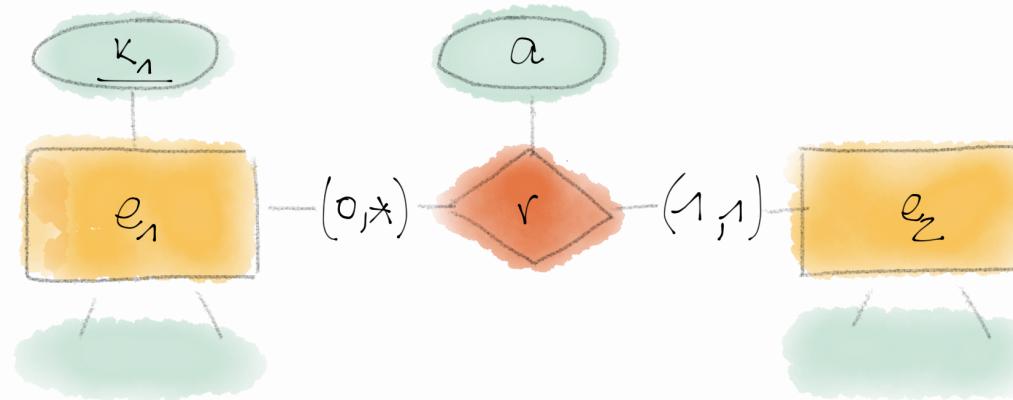


1. Assume table $t_i = \text{plural}(e_i)$ translates e_i .
2. In t_2 , add column r as a **foreign key** referencing table t_1 (**implements relationship** r).
3. Replace key of t_2 with **composite primary key** (r, k_2) .
4. Implement existence dependency via an **ON DELETE CASCADE** action on foreign key r :

```
ALTER TABLE t2 ADD COLUMN r <type of k1> NOT NULL;
ALTER TABLE t2 DROP CONSTRAINT t2_pkey;    -- ← replace primary key
ALTER TABLE t2 ADD PRIMARY KEY (r, k2);    -- ← of table t2
ALTER TABLE t2 ADD FOREIGN KEY (r) REFERENCES t1(k1) ON DELETE CASCADE;
```

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #3: One-to-Many Relationship Types



1. Assume table $t_i = \text{plural}(e_i)$ translates e_i .
2. Implement r and a : add column r to t_2 and establish as a foreign key referencing table t_1 (no dependency, thus no **ON DELETE** action). Add column a to t_2 :

```
-- establish foreign key to implement relationship r
ALTER TABLE t2 ADD COLUMN r <type of k1> NOT NULL; -- ← realize (1,...) card
ALTER TABLE t2 ADD COLUMN a <type of a> NOT NULL; -- ←
ALTER TABLE t2 ADD FOREIGN KEY (r) REFERENCES t1(k1);
```

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #3: One-to-Many Relationship Types (Cardinalities)

1. $-(0,*)-\diamond-(1,1)-$

The key value of a given e_1 entity does *not* necessarily appear in column r of t_2 .

2. $-(0,*)-\diamond-(1,1)-$

The key value of a given e_1 entity may appear *multiple times* in column r of t_2 .

3. $-(0,*)-\diamond-(1,1)-$

Value of column r in t_2 is **NOT NULL** (remove **NOT NULL** for $-(0,*)-\diamond-(0,1)-$).

4. $-(0,*)-\diamond-(1,1)-$

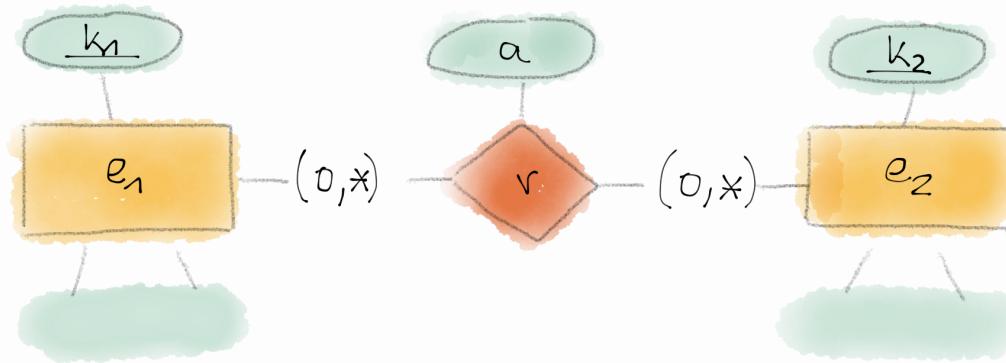
Column r has exactly *one value* in each row of t_2 .

⚠ **But:** Relational **NOT NULL**, key, and foreign key constraints **cannot implement general ER cardinality constraints**. Unimplementable:

- $-(m_1,*)-\diamond-(_,1)-$ with $m_1 > 0$ and $-(0,n_1)-\diamond-(_,1)-$ with $n_1 > 0$, $n_1 \neq *$

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #3b: Many-to-Many Relationship Types



1. Assume table $t_i = \text{plural}(e_i)$ translates e_i .
2. Implement r and a : **CREATE TABLE** r with columns k_1 , k_2 , a to represent the relationship. Establish (k_1, k_2) as **composite key** for table r . In table r , establish k_i as foreign key referencing table t_i with **ON DELETE** actions.

```
CREATE TABLE r (k1 <type of k1>, k2 <type of k2>, a <type of a>);  
ALTER TABLE r ADD PRIMARY KEY (k1, k2);  
ALTER TABLE r ADD FOREIGN KEY (k1) REFERENCES t1(k1) ON DELETE CASCADE;  
ALTER TABLE r ADD FOREIGN KEY (k2) REFERENCES t2(k2) ON DELETE CASCADE;
```

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #3b: Many-to-Many Relationship Types

- Quiz:

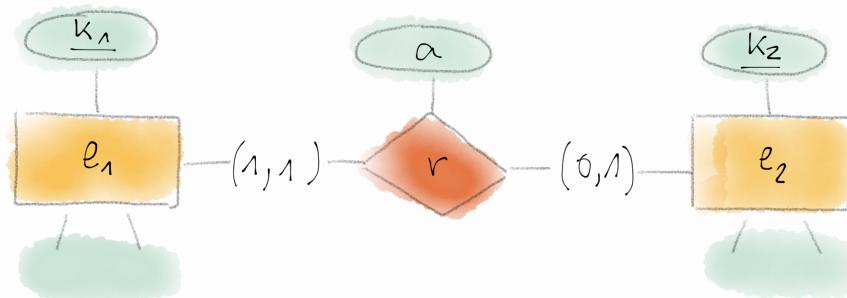
1. Explain the choice of *composite primary key* (k_1, k_2) for table r .
2. What kind of mini-world is implemented if we establish column k_1 (or k_2) as the *only* key column in table r ? Is the result useful?

- Notes on many-to-many relationship cardinalities:

- Cardinalities other than $-(0,*)-\diamond-(0,*)-$ cannot be enforced by the relational model. In particular, $-(m,*)-\diamond-(_,*)-$ with $m > 0$ cannot be translated faithfully.
- If such scenarios are important in mini-world modelling, the cardinality constraint needs to be checked by
 1. the application program or
 2. a general SQL constraint mechanism (non-relational, in the strict sense).

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #3c: One-to-One Relationship Types



1. Assume table $t_i = \text{plural}(e_i)$ translates e_i .
2. Two options: host **implementation of r and a** in t_1 or t_2 . Here: choose t_1 since **every e_1 entity participates in r -relationship** $-(1,1)-\diamond\dots$: Add column r to t_1 and establish as a foreign key referencing table t_2 (no dependency, thus no **ON DELETE** action). Add column a to t_1 .

```
ALTER TABLE t1 ADD COLUMN r <type of k2> NOT NULL;
ALTER TABLE t1 ADD FOREIGN KEY (r) REFERENCES t2(k2);
ALTER TABLE t1 ADD UNIQUE (r); -- A r becomes candidate key in t1
ALTER TABLE t1 ADD COLUMN a <type of a> NOT NULL;
```

MAPPING EER TO RELATIONAL TARGET (SQL)

Step #3c: One-to-One Relationship Types (Cardinalities)

- Consider how the relational translation implements the ER cardinality constraints:

1. $-(1,1)-\diamond-(0,1)-$:

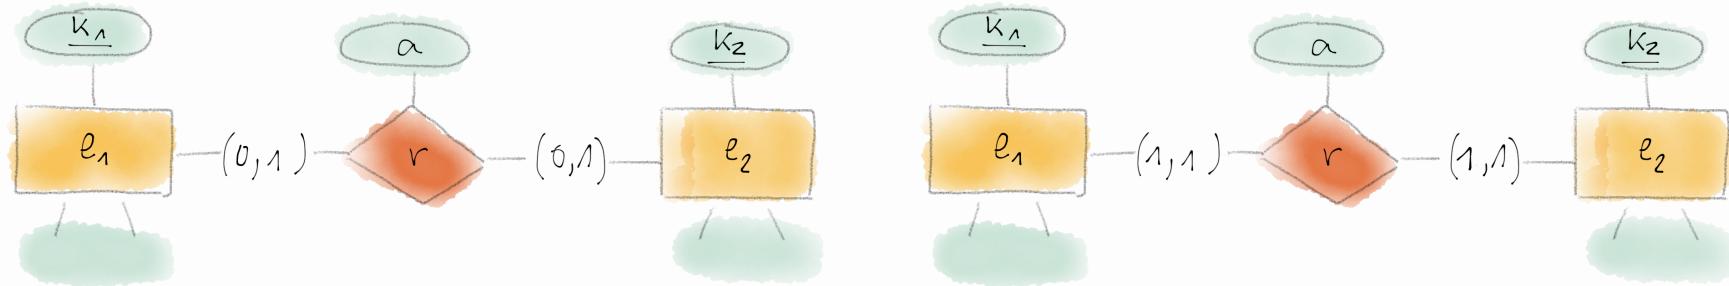
2. $-(1,1)-\diamond-(0,1)-$:

3. $-(1,1)-\diamond-(0,1)-$:

4. $-(1,1)-\diamond-(0,1)-$:

MAPPING EER TO RELATIONAL TARGET (SQL)

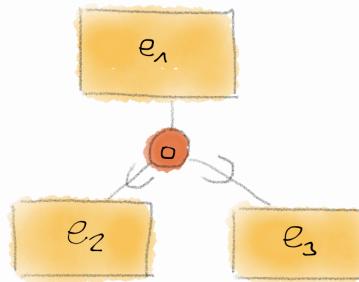
Step #3c: One-to-One Relationship Types (Alternatives)



- $-(0,1)-\diamond-(0,1)-$: Two options:
 1. Drop **NOT NULL** constraint on column r in table t_1 .
 2. Represent r (and a) in a **separate table r** with two foreign keys k_i referencing t_i . Both k_i are candidate keys in table r .
- $-(1,1)-\diamond-(1,1)-$:
Merge tables t_1 and t_2 and column a into single table t . No **NULL** values allowed to enforce minimum cardinalities of **1**. Both k_i are candidate keys in table t .

MAPPING EER TO RELATIONAL TARGET (SQL)

Translating Inheritance



1. Assume table $t_i = \text{plural}(e_i)$ will hold the translation of e_i .
2. Implement partial, overlapping inheritance: **CREATE TABLEs** t_2 and t_3 using SQL's **INHERITS** declaration to indicate that both inherit from table t_1 .

```
CREATE TABLE t1 (...);  
CREATE TABLE t2 (...) INHERITS (t1);  
CREATE TABLE t3 (...) INHERITS (t1);
```

3. The implementation of **total** or **disjoint** inheritance requires a bit more SQL machinery (e.g., PostgreSQL's **CREATE RULE** mechanism, see below).

SQL: TABLE INHERITANCE VIA INHERITS

INHERITS (PostgreSQL)

If **subtable** s **INHERITS** from tables t_1, \dots, t_n ,

```
CREATE TABLE s(...) [ INHERITS (t1 [, ...]) ]
```

1. subtable s features all columns of tables t_i and those declared locally for s , i.e., $sch(t_i) \subseteq sch(s)$,
2. rows inserted/updated/deleted in s will also affect all tables t_i (but not vice versa), i.e., $inst(s)|sch(t_i) \subseteq inst(t_i)$ ($|C \equiv$ restricted to the columns in set C).

- ⚠ In PostgreSQL, the subtable does *not* inherit key, foreign key, or **UNIQUE** constraints of its supertables. Re-declare these for the subtable.
- Note: **TABLE** t_i produces $inst(t_i)$. **TABLE ONLY** t_i produces $inst(t_i) \setminus (inst(s)|sch(t_i))$.

THE COMPLETE LEGO DATABASE ER MODEL

